

Predict Customer Churn in Streaming Service Model

Toan Nguyen

06 July 2025

Contents

Introduction	1
Library setup (Assume all packages are pre-installed)	2
Data loading	2
Building decision tree models	4
Parameter Tuning	5
Check and verify which values are more suitable for each parameter	6
Tree 1: Using default parameters	6
Tree 2: Limit the maximum depth of tree	7
Tree 3: Complex tree with many splits	9
Tree 4: Pruned tree with the best complexity parameter	10
Model Performance Summary	11
Performance Metric Definitions	15
Model comparison	15
Recommendation	16

Introduction

In the context of predicting customer churn for our subscription-based digital streaming service, a decision tree model was chosen for its ability to balance accuracy with interpretability. Our dataset includes a diverse mix of variables such as subscription type, preferred device, payment frequency, customer demographics, and engagement metrics like total watch time and number of support tickets. A decision tree model is an appropriate initial choice given its ability to handle both numerical and categorical variables with minimal preprocessing. However, its performance should be carefully reassessed alongside other predictive models to ensure the most accurate and reliable approach is selected for forecasting customer churn.

More importantly, the model generates easily understandable rules that allow us to pinpoint the conditions under which customers are most likely to cancel like low engagement combined with basic-tier subscriptions or high support ticket volumes. This level of transparency is essential for business decision-making, as it enables marketing, product, and support teams to take targeted action based on clear churn indicators. While more advanced ensemble models may offer marginal improvements in predictive accuracy, the decision tree provides a solid foundation for identifying actionable patterns and building a churn mitigation strategy grounded in data-driven insights.

Library setup (Assume all packages are pre-installed)

```
library(readr)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(rpart)
library(rattle)
```

```
## Loading required package: tibble
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.
## Version 5.5.1 Copyright (c) 2006-2021 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
```

```
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      cov, smooth, var
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

Data loading

Load the csv file and take a look at first 5 rows

```
streaming_data <- read_csv('streaming_service_cleaned.csv')
```

```
## Rows: 2375 Columns: 19
## -- Column specification -----
## Delimiter: ","
## chr  (10): Name, Gender, SubscriptionType, PaymentMethod, Email, Country, Re...
## dbl  (6): CustomerID, Age, MonthlyFee, TotalWatchTime, NumSupportTickets, L...
## lgl  (1): Cancelled
## dtm   (2): JoinDate, LastLoginDate
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
head(streaming_data)
```

```
## # A tibble: 6 x 19
##   CustomerID Name      Age Gender JoinDate      SubscriptionType MonthlyFee
##   <dbl> <chr>   <dbl> <chr>   <dtm>         <chr>             <dbl>
## 1         1 David~    19 Male  2022-06-29 00:00:00 Standard           16.4
## 2         2 Emily~    53 Non-b~ 2021-12-15 00:00:00 Standard           15.9
## 3         3 Alex ~    75 Male  2023-09-11 00:00:00 Premium            18.7
## 4         4 Emily~    58 Female 2019-06-15 00:00:00 Basic              10.7
## 5         5 Jane ~    27 Male  2022-03-17 00:00:00 Basic               8.98
## 6         6 David~    58 Male  2023-04-01 00:00:00 Basic               9.72
## # i 12 more variables: PaymentMethod <chr>, TotalWatchTime <dbl>,
## #   Cancelled <lgl>, Email <chr>, Country <chr>, Region <chr>,
## #   LastLoginDate <dtm>, NumSupportTickets <dbl>, LoyaltyPoints <dbl>,
## #   ReferralSource <chr>, PreferredDevice <chr>, PaymentFrequency <chr>
```

```
# Check number of rows
```

```
nrow(streaming_data)
```

```
## [1] 2375
```

Select features for predicting (based on previous data exploratory performance) with target variable “Cancelled”

```
# Filter to use only the specified variables
streaming_data_filtered <- streaming_data %>% select(
  Region,
  Country,
  PreferredDevice,
  Gender,
  SubscriptionType,
  PaymentFrequency,
  PaymentMethod,
  Age,
  NumSupportTickets,
  ReferralSource,
  Cancelled
)
```

Make sure all variables have appropriate types

```
# Convert variables to appropriate types
streaming_data_filtered$Region <- as.factor(streaming_data_filtered$Region)
streaming_data_filtered$Country <- as.factor(streaming_data_filtered$Country)
streaming_data_filtered$PreferredDevice <-
  ↳ as.factor(streaming_data_filtered$PreferredDevice)
streaming_data_filtered$Gender <- as.factor(streaming_data_filtered$Gender)
streaming_data_filtered$SubscriptionType <-
  ↳ as.factor(streaming_data_filtered$SubscriptionType)
streaming_data_filtered$PaymentFrequency <-
  ↳ as.factor(streaming_data_filtered$PaymentFrequency)
streaming_data_filtered$PaymentMethod <- as.factor(streaming_data_filtered$PaymentMethod)
streaming_data_filtered$Age <- as.integer(streaming_data_filtered$Age)
streaming_data_filtered$NumSupportTickets <-
  ↳ as.integer(streaming_data_filtered$NumSupportTickets)
streaming_data_filtered$ReferralSource <-
  ↳ as.factor(streaming_data_filtered$ReferralSource)
streaming_data_filtered$Cancelled <- as.factor(streaming_data_filtered$Cancelled)

# Basic streaming_data_filtered summary
summary(streaming_data_filtered)
```

```
##      Region      Country PreferredDevice      Gender
## North   :260  Australia:385 Desktop: 684   Female   :1103
## Region9 :254   Canada   :380 Mobile :1455   Male     :1033
## Region7 :249   France   :368 Tablet  : 236   Non-binary : 114
## Region10:244   Germany :414                Prefer not to say: 125
## South   :240    UK       :408
## Central :236    USA      :420
## (Other) :892
## SubscriptionType PaymentFrequency      PaymentMethod      Age
## Basic   :967   Annual   : 243 Bank Transfer: 246   Min.    :18.00
## Premium :450   Monthly  :1662 Credit Card  :1134   1st Qu.:34.00
## Standard:958   Quarterly: 470 Debit Card   : 513   Median  :49.00
##                                     Paypal       : 482   Mean    :49.33
##                                     3rd Qu.:64.00
##                                     Max.    :80.00
##
## NumSupportTickets      ReferralSource Cancelled
## Min.    :0.000      Ad           :481   FALSE:1692
## 1st Qu.:1.000      Friend        :522   TRUE : 683
## Median  :1.000      OrganicSearch:238
## Mean    :1.689      SocialMedia  :447
## 3rd Qu.:2.000      Website      :687
## Max.    :6.000
##
```

Building decision tree models

As seen in the summary, the target variable “Cancelled” consist of total 2375 samples, with 29% of TRUE and 71% of FALSE.

This dataset is slightly imbalanced, therefore a 70/30 split with stratify can help with evaluating the model more reliably while keeping the class distribution.

```
set.seed(42) # for reproducibility

# Create stratified split (70% training, 30% testing)
train_index <- createDataPartition(streaming_data_filtered$Cancelled,
                                   p = 0.7,
                                   list = FALSE)

train_data <- streaming_data_filtered[train_index, ]
test_data  <- streaming_data_filtered[-train_index, ]

# Check class distribution
print("Train data distribution:")
```

```
## [1] "Train data distribution:"
```

```
prop.table(table(train_data$Cancelled))
```

```
##
##      FALSE      TRUE
## 0.7121394 0.2878606
```

```
print("Test data distribution:")
```

```
## [1] "Test data distribution:"
```

```
prop.table(table(test_data$Cancelled))
```

```
##
##      FALSE      TRUE
## 0.7130802 0.2869198
```

The next step is to fit the decision tree.

Parameter Tuning

There are lots of parameters, however these are the most relevant:

Parameters:

Parameter	Description
maxdepth	Maximum depth of any node of the final tree (limits complexity).
minsplit	Minimum number of observations required to attempt a split.
cp	Complexity parameter: smaller values allow more splits. Used for pruning and controlling overfitting.
xval	number of cross-validations.

Check and verify which values are more suitable for each parameter

```
for (cp_val in c(0.01, 0.005, 0.001)) {  
  for (minsplit_val in c(10, 15, 20)) {  
    model <- rpart(Cancelled ~ ., data = train_data, method = "class",  
                  control = rpart.control(cp = cp_val, minsplit = minsplit_val, maxdepth  
                    = 10))  
    pred <- predict(model, newdata = test_data, type = "class")  
    acc <- mean(pred == test_data$Cancelled)  
    cat("cp =", cp_val, "minsplit =", minsplit_val, "Accuracy =", acc, "\n")  
  }  
}
```

```
## cp = 0.01 minsplit = 10 Accuracy = 0.7130802  
## cp = 0.01 minsplit = 15 Accuracy = 0.7130802  
## cp = 0.01 minsplit = 20 Accuracy = 0.7130802  
## cp = 0.005 minsplit = 10 Accuracy = 0.7130802  
## cp = 0.005 minsplit = 15 Accuracy = 0.7130802  
## cp = 0.005 minsplit = 20 Accuracy = 0.7130802  
## cp = 0.001 minsplit = 10 Accuracy = 0.6202532  
## cp = 0.001 minsplit = 15 Accuracy = 0.6469761  
## cp = 0.001 minsplit = 20 Accuracy = 0.6540084
```

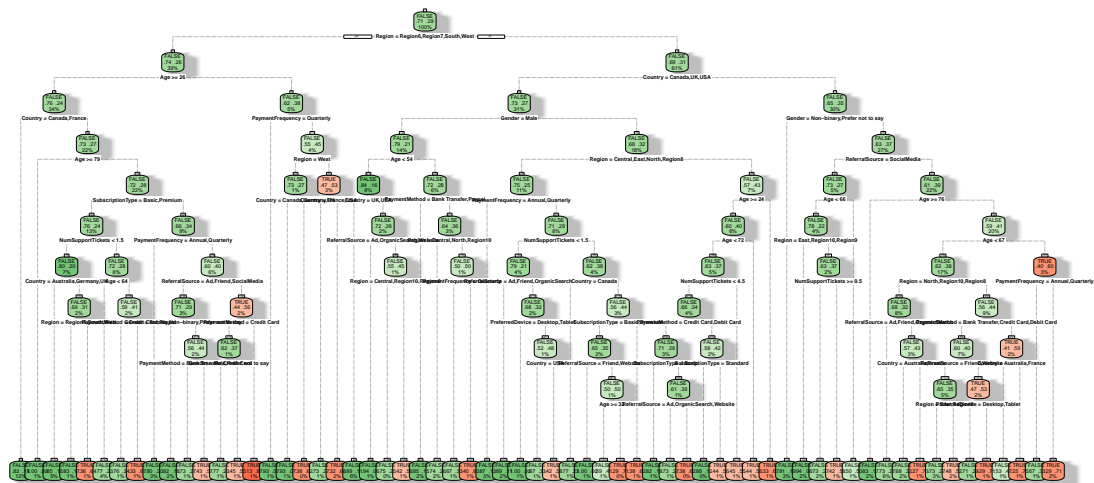
Models with $cp = 0.01$ and $cp = 0.005$, regardless of the $minsplit$ value, all resulted in an accuracy of 71.3%. However, this matches the proportion of the majority class ($Cancelled = FALSE$), indicating that the trees were over-pruned and likely made no meaningful splits - simply defaulting to predicting the most common outcome. In contrast, lowering the complexity parameter to $cp = 0.001$ allowed the tree to start forming actual branches. Among these, the model with $minsplit = 20$ achieved the best accuracy at 65.4%, which, although lower than the previous models, reflects a genuine attempt to classify both TRUE and FALSE cases of cancellation. This suggests that it is the first model to move beyond the baseline and begin learning from the features in the data.

Tree 1: Using default parameters

Using $cp = 0.001$, $minsplit = 20$, $maxdepth = 10$

```
tree1 <- rpart(  
  Cancelled ~ .,  
  data = train_data,  
  method = "class",  
  control = rpart.control(  
    maxdepth = 10,  
    minsplit = 20,  
    cp = 0.001  
  )  
)  
fancyRpartPlot(  
  tree1,  
  palettes = c("Greens", "Reds"),  
  sub = "Default Parameters"  
)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Default Parameters

Figure 1: Decision Tree 1

Export the tree model as an image so we can view clearer.

```
# Save as large PNG
png("tree_1.png", width = 5000, height = 1800, res = 200)
fancyRpartPlot(tree1, palettes = c("Greens", "Reds"), sub = "Pruned Tree", cex = 0.6)
dev.off()
```

```
## pdf
## 2
```

Tree 2: Limit the maximum depth of tree

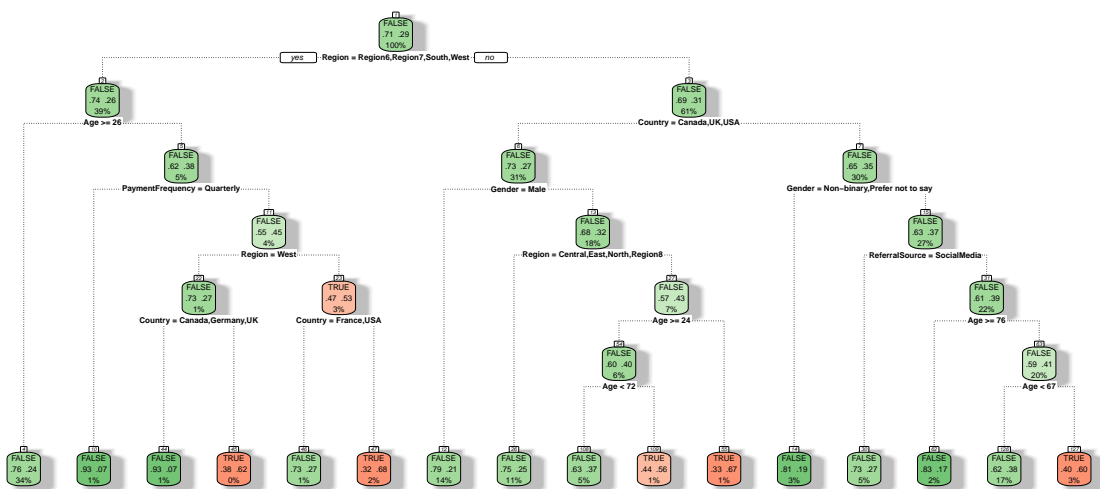
Using $cp = 0.001$, $minsplit = 20$, $maxdepth = 6$

```
tree2 <- rpart(
  Cancelled ~ .,
  data = train_data,
  method = "class",
```

```

control = rpart.control(
  maxdepth = 6,
  minsplit = 20,
  cp = 0.001
)
fancyRpartPlot(
  tree2,
  palettes = c("Greens", "Reds"),
  sub = "Default Parameters"
)

```



Default Parameters

Figure 2: Decision Tree 2

Export the tree model as an image so we can view clearer.

```

# Save as large PNG
png("tree_2.png", width = 5000, height = 1800, res = 200)
fancyRpartPlot(tree1, palettes = c("Greens", "Reds"), sub = "Pruned Tree", cex = 0.6)
dev.off()

```

```

## pdf
## 2

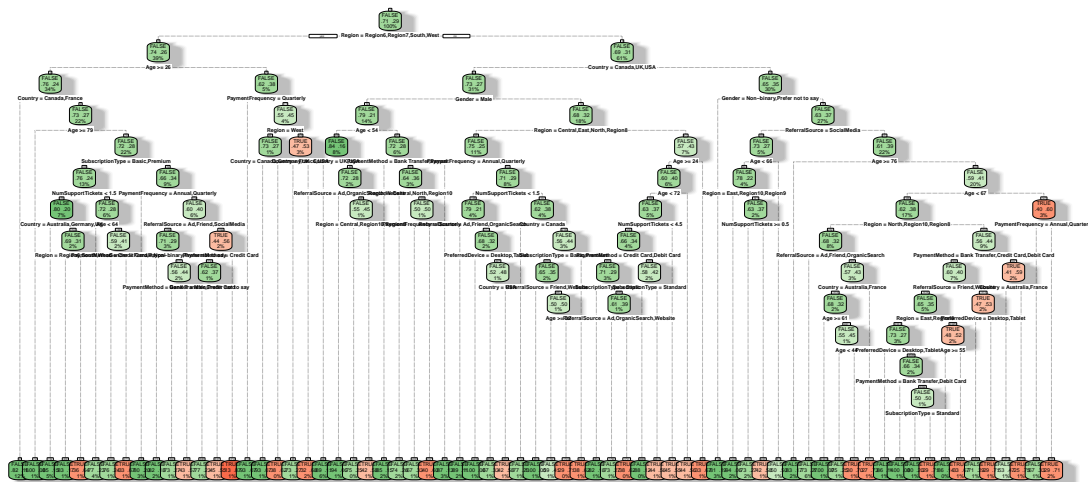
```


Tree 3: Complex tree with many splits

Using $cp = 0.001$, $minsplit = 20$, $maxdepth = 30$

```
tree3 <- rpart(  
  Cancelled ~ .,  
  data = train_data,  
  method = "class",  
  control = rpart.control(  
    maxdepth = 30,  
    minsplit = 20,  
    cp = 0.001  
  )  
)  
fancyRpartPlot(  
  tree3,  
  palettes = c("Greens", "Reds"),  
  sub = "Default Parameters"  
)
```

Warning: labs do not fit even at cex 0.15, there may be some overplotting



Default Parameters

Figure 3: Decision Tree 3

Export the tree model as an image so we can view clearer.

```
# Save as large PNG
png("tree_3.png", width = 5000, height = 1800, res = 200)
fancyRpartPlot(tree1, palettes = c("Greens", "Reds"), sub = "Pruned Tree", cex = 0.6)
dev.off()
```

```
## pdf
## 2
```

Tree 4: Pruned tree with the best complexity parameter

Using $cp = 0$

Expected impact: Simpler tree with balance performances

```
# Build the full decision tree with the specified complexity parameter
c.tree.full <- rpart(
  Cancelled ~ .,
  data = train_data,
  method = "class",
  control = rpart.control(cp = 0)
)

# Prune the tree using the specified CP value
tree4 <- prune(c.tree.full, cp = 0)

# Visualize the pruned tree
fancyRpartPlot(tree4, palettes = c("Greens", "Reds"), sub = "")
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```

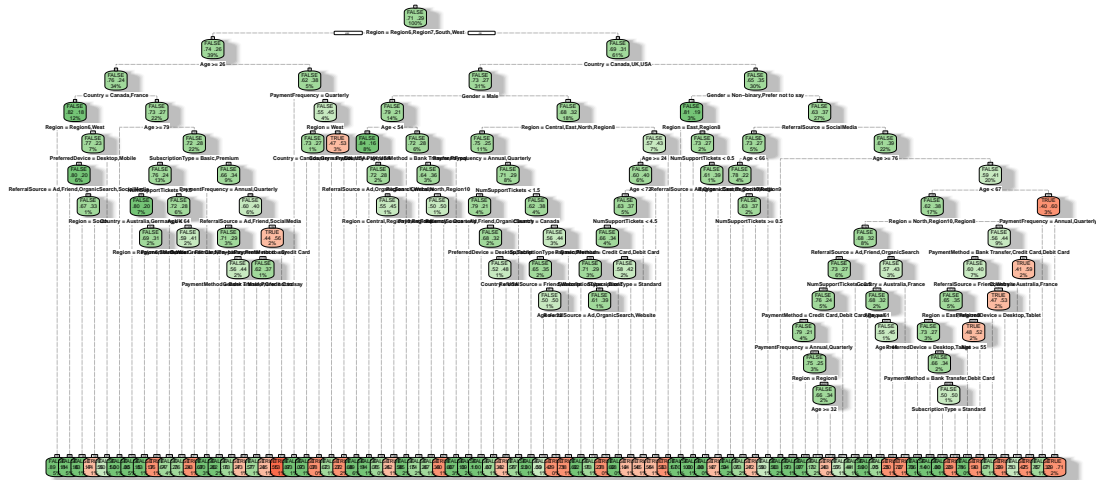


Figure 4: Decision Tree 4

Export the tree model as an image so we can view clearer.

```
# Save as large PNG
png("tree_4.png", width = 5000, height = 1800, res = 200)
fancyRpartPlot(tree1, palettes = c("Greens", "Reds"), sub = "Pruned Tree", cex = 0.6)
dev.off()
```

```
## pdf
## 2
```

Model Performance Summary

```
# Ensure factor levels are consistent
true_labels <- factor(test_data$Cancelled, levels = c("FALSE", "TRUE"))
```

```
# List to store ROC objects
roc_list <- list()
```

```
# Loop through tree models 1-4
for (i in 1:4) {
```

```

model <- get(paste0("tree", i))

# Class predictions
pred_class <- predict(model, test_data, type = "class")
pred_class <- factor(pred_class, levels = c("FALSE", "TRUE"))

# Probabilistic predictions for AUC
pred_prob <- predict(model, test_data, type = "prob")[, "TRUE"]

# Confusion matrix with accuracy CIs
cm <- confusionMatrix(
  data      = pred_class,
  reference = true_labels,
  positive  = "TRUE"
)

# Extract key metrics
accuracy    <- cm$overall["Accuracy"]
acc_lower   <- cm$overall["AccuracyLower"]
acc_upper   <- cm$overall["AccuracyUpper"]
precision   <- cm$byClass["Precision"]
recall      <- cm$byClass["Recall"]
f1          <- cm$byClass["F1"]

# ROC and AUC
roc_obj     <- pROC::roc(response = true_labels, predictor = pred_prob)
roc_list[[paste0("Tree ", i)]] <- roc_obj
auc_value   <- pROC::auc(roc_obj)

# Report
cat("\n===== \n")
cat("Tree", i, "Performance: \n")
cat("===== \n")
print(cm$table)
cat(sprintf(
  "Accuracy      : %.3f (95%% CI: %.3f-%.3f) \n",
  accuracy, acc_lower, acc_upper
))
cat(sprintf("Misclass. err : %.3f \n", 1 - accuracy))
cat(sprintf("Precision   : %.3f \n", precision))
cat(sprintf("Recall      : %.3f \n", recall))
cat(sprintf("F1-Score    : %.3f \n", f1))
cat(sprintf("AUC         : %.3f \n", auc_value))
}

```

```
## Setting levels: control = FALSE, case = TRUE
```

```
## Setting direction: controls < cases
```

```
##
## =====
## Tree 1 Performance:
## =====
```

```
##           Reference
## Prediction FALSE TRUE
##      FALSE   422  161
##      TRUE     85   43
## Accuracy      : 0.654 (95% CI: 0.618-0.689)
## Misclass. err : 0.346
## Precision     : 0.336
## Recall        : 0.211
## F1-Score      : 0.259
## AUC           : 0.546
```

```
## Setting levels: control = FALSE, case = TRUE
## Setting direction: controls < cases
```

```
##
## =====
## Tree 2 Performance:
## =====
##           Reference
## Prediction FALSE TRUE
##      FALSE   478  193
##      TRUE     29   11
## Accuracy      : 0.688 (95% CI: 0.652-0.722)
## Misclass. err : 0.312
## Precision     : 0.275
## Recall        : 0.054
## F1-Score      : 0.090
## AUC           : 0.546
```

```
## Setting levels: control = FALSE, case = TRUE
## Setting direction: controls < cases
```

```
##
## =====
## Tree 3 Performance:
## =====
##           Reference
## Prediction FALSE TRUE
##      FALSE   418  161
##      TRUE     89   43
## Accuracy      : 0.648 (95% CI: 0.612-0.684)
## Misclass. err : 0.352
## Precision     : 0.326
## Recall        : 0.211
## F1-Score      : 0.256
## AUC           : 0.534
```

```
## Setting levels: control = FALSE, case = TRUE
## Setting direction: controls < cases
```

```
##
## =====
```

```
## Tree 4 Performance:
## =====
##           Reference
## Prediction FALSE TRUE
##      FALSE   410  158
##      TRUE     97   46
## Accuracy      : 0.641 (95% CI: 0.605-0.677)
## Misclass. err : 0.359
## Precision     : 0.322
## Recall        : 0.225
## F1-Score      : 0.265
## AUC           : 0.502

# Set up the plotting area with 2x2 grid
par(mfrow = c(2, 2))

# Color palette for the plots
cols <- c("#1f77b4", "#ff7f0e", "#2ca02c", "#d62728")

# Loop through the ROC list and plot each
for (i in 1:4) {
  plot(
    roc_list[[i]],
    col = cols[i],
    main = paste("Tree Model", i),
    lwd = 2,
    xlim = c(1, 0),
    ylim = c(0, 1)
  )

  # Calculate and display the AUC value
  auc_value <- auc(roc_list[[i]])
  text(
    x = 0.4, y = 0.15,
    labels = paste("AUC:", round(auc_value, 3)),
    col = cols[i],
    cex = 1.2,
    font = 2
  )
}
}
```

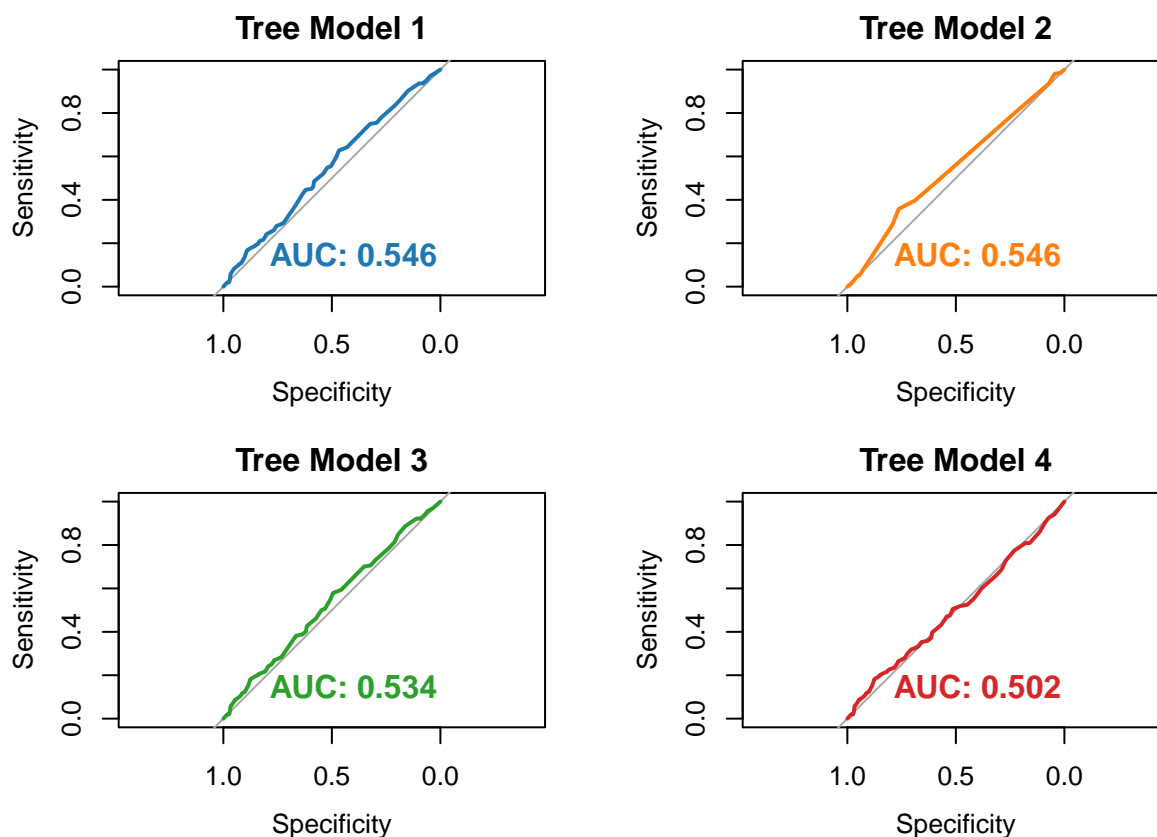


Figure 5: AUC comparison

Performance Metric Definitions

Performance Metric	Definition
Accuracy	How many predictions were correct overall.
Precision	Out of all predicted churners, how many actually churned.
Recall	Out of all actual churners, how many we correctly predicted.
F1-Score	A balance between precision and recall.
AUC	How well the model separates churners from non-churners.

Model comparison

In evaluating the decision tree models for predicting customer churn, we compared performance across multiple key metrics, including accuracy, precision, recall, F1-score, and AUC, all of which are crucial in understanding how well each model can help the business proactively retain customers.

Tree 2 achieved the highest accuracy (68.8%), with a 95% Confidence Interval (CI) of [65.2%, 72.2%], suggesting that it correctly predicts churn and non-churn outcomes more often than the other models. However, accuracy alone isn't

enough when dealing with customer churn, because a high accuracy can be misleading if the model fails to identify the customers most likely to churn. Tree 2 has a very low recall (5.4%), which means it misses a significant portion of actual churned customers. In the context of customer churn, recall is critical because it tells us how many of the customers at risk of leaving we are able to identify. Missing churned customers reduces our ability to take action and retain them. Thus, despite Tree 2's high accuracy, its low recall makes it unsuitable for proactively targeting customers who are likely to churn.

Tree 1, with an accuracy of 65.4% and a 95% CI of [61.8%, 68.9%], performs better than Tree 2 in terms of recall (21.1%). This indicates that it captures more of the churned customers compared to Tree 2, though still not enough. Precision (33.6%) in Tree 1 is also higher than Tree 2, meaning it is less likely to incorrectly flag non-churned customers as churned. However, its F1-score (0.259)—which balances precision and recall—shows that the model is still not good enough at identifying churned customers while avoiding false positives. The trade-off between precision and recall in Tree 1 means that it does better at reducing errors but still misses a significant portion of the customers we want to retain.

Tree 3, with an accuracy of 64.8% and a 95% CI of [61.2%, 68.4%], performs similarly to Tree 1 in terms of recall (21.1%) but slightly outperforms Tree 1 in precision (32.6%). This means that Tree 3 is more accurate in predicting non-churned customers, but still fails to predict a sufficient number of churned customers. While it improves slightly over Tree 1, its F1-score (0.256) is still low, indicating that it's not the best at balancing precision and recall for churn prediction.

In contrast, Tree 4 has a lower accuracy (64.1%), with a 95% CI of [60.5%, 67.7%], but its strength lies in its recall (22.5%), which is the highest among all the models. This is extremely important for predicting customer churn because recall tells us how many actual churned customers the model successfully identifies. A high recall means that the business can act quickly to prevent customers from leaving, which is the ultimate goal when trying to reduce churn. Although Tree 4 has a slightly lower accuracy and the lowest AUC (0.502), its recall is the best, allowing the company to focus on the customers who are most likely to leave and offer interventions to retain them. Tree 4 also achieves a solid F1-score (0.265), striking a balance between reducing false positives and capturing churned customers.

Considering the business objective of reducing customer churn, Tree 4 is the most effective model despite having a slightly lower accuracy compared to Tree 2. Its higher recall is critical in identifying customers who are at risk of leaving, enabling the business to act early and improve retention efforts. Tree 4 captures more churned customers than the other models, making it the best choice for proactively targeting and retaining at-risk customers. Therefore, Tree 4 should be selected as the primary model for churn prediction in customer retention strategies.

Recommendation

The decision tree model 4 highlights several key variables influencing customer churn, starting with geographic region (Please refer to tree_4.png). Customers from areas such as France, Australia, and the United States tend to cancel more frequently, suggesting regional preferences or competitive pressure may affect retention. In response, the business should design localized campaigns tailored to each market, including language-specific support, regionally relevant content, and targeted promotional offers. Payment behavior also plays a critical role. Customers billed annually or quarterly, especially those using bank transfers or debit cards, are more likely to churn. To address this, the company should offer greater billing flexibility, provide reminders before renewal dates, and encourage smoother payment methods such as credit cards or digital wallets. Subscription type is another important factor. Users on basic or premium plans show higher churn, which may indicate dissatisfaction with features or value. A review of these plans could lead to adjustments such as bundling more content into premium or enhancing the basic tier to better match customer expectations.

Referral source also affects loyalty. Customers acquired through advertisements or search engines are more likely to leave, while those referred by friends or loyalty programs tend to stay longer. This insight supports expanding referral initiatives by offering dual-sided incentives, simplifying referral tracking, and promoting these programs within the platform. Age is another predictor, with younger customers under 26 showing a higher likelihood of cancellation. The business can support this segment through onboarding tutorials, personalized content recommendations, and early engagement strategies to build loyalty. Lastly, the number of support tickets submitted by a customer serves as a valuable signal. Very low ticket activity may suggest disengagement, while high volumes could indicate unresolved

frustration. The company should monitor these patterns closely and create triggers for proactive outreach, such as check-in emails or priority support offers, to intervene before churn occurs.

While Model 4 provides actionable insights, its performance must be weighed carefully. It offers a recall of 22.5 percent, meaning it can identify many customers at risk of leaving, but its overall accuracy and precision are moderate. This could lead to overestimating churn, resulting in unnecessary retention costs. Therefore, it is essential to compare this model with others. Logistic regression is easy to interpret but may miss complex relationships. Models such as random forests, gradient boosting, and XGBoost offer greater accuracy and stronger predictive performance but require more resources and may sacrifice some transparency. Comparing these models using consistent metrics—accuracy, precision, recall, F1-score, and AUC—will help determine which approach offers the most reliable and useful results. A balanced model will allow the business to identify churn risk with confidence, focus retention efforts where they matter most, and improve long-term customer value.